

INTERNSHIP PROJECT REPORT

VisionSpec QC

AI-Powered PCB Visual Inspection System

Submitted by

Bharath Krishnan S

sbkcuddalore@gmail.com

Organization: Zaalima Development

Duration: 4 Weeks | June 2026

1. Internship & Project Details

Intern Name	Bharath Krishnan S
Email	sbkcuddalore@gmail.com
Organization	Zaalima Development
Project Title	VisionSpec QC — AI-Powered PCB Visual Inspection System
Domain	Artificial Intelligence / Computer Vision
Duration	4 Weeks
Report Date	June 2026
Status	Completed — running end-to-end

2. Project Overview (Abstract)

VisionSpec QC is an end-to-end, AI-powered quality-control application that automates the visual inspection of Printed Circuit Boards (PCBs). The system classifies each board as **Pass** or **Defect** and applies Explainable AI (Grad-CAM) to highlight the exact regions responsible for a defect verdict, removing the “black-box” limitation of deep neural networks.

The platform is built as a modular full-stack system: a deep-learning model (MobileNetV2 transfer learning) served through a FastAPI backend, persisted to a SQLite database, and presented through a real-time React + TypeScript dashboard supporting image upload, live webcam inference, analytics, and inspection history.

3. Objectives

- Build an automated PCB defect-classification model using transfer learning.
- Provide Explainable AI (Grad-CAM heatmaps) to justify each prediction.
- Expose real-time inference through production-grade REST APIs (FastAPI).
- Deliver an interactive dashboard with live detection, analytics, and history.
- Persist every prediction for auditing and a Human-in-the-Loop feedback workflow.

4. Technologies Used

Layer	Technology
AI / Machine Learning	TensorFlow 2, Keras, MobileNetV2 (Transfer Learning), Grad-CAM
Computer Vision	OpenCV, NumPy, Pillow
Backend	FastAPI, Uvicorn, SQLAlchemy, SQLite
Frontend	React, TypeScript, Vite, Chart.js, Axios

Layer	Technology
Tooling	Python venv, Git, psutil, Matplotlib (reports)

5. Weekly Progress Breakdown

Week 1 — Foundation, Dataset & Project Scaffolding

Goal: Establish the development environment, raw data, and skeleton of all three layers.

- Set up Python virtual environment and installed the core AI/backend stack; initialized the React + TypeScript app via Vite.
- Designed the modular backend (routes, services, models, schemas) and frontend (pages, components, services) structure.
- Built a procedural **Synthetic PCB Dataset Generator** producing 1,000 images (500 Pass / 500 Defect) at 224×224, with injected defects (scratches, burns, oxidation, missing traces).
- Configured Keras ImageDataGenerator augmentation and created routing for the core pages.

Week 2 — Core Detection Model & Prediction API

Goal: Train the first working classifier and connect it to an upload API and UI.

- Built a MobileNetV2 transfer-learning model: GlobalAveragePooling → Dense(256) → Dense(128) → Dense(1, Sigmoid).
- Two-phase training: 10 epochs frozen backbone, then fine-tuned top layers for 10 more epochs (early stop at epoch 17); exported pcb_model.h5.
- Developed the POST /api/predict endpoint (validation, preprocessing, inference, response with label, confidence, latency).
- Built the Defect Analysis page with drag-and-drop upload wired to the prediction API.

Week 3 — Explainable AI (Grad-CAM) & Live Video Pipeline

Goal: Add visual explainability and a live webcam inference path.

- Implemented Grad-CAM against MobileNetV2's final conv layer (out_relu) to generate defect heatmaps.
- Encoded heatmaps as base64 in the prediction response and rendered them as an overlay in the UI.
- Built the POST /api/webcam/predict endpoint using OpenCV for frame-by-frame inference.
- Created the Live Detection page that streams the browser webcam feed with real-time results and FPS monitoring.

Week 4 — Analytics, Logging, Dashboard & Integration

Goal: Tie everything together with persistence, analytics, evaluation, and polish.

- Added SQLite prediction logging and the GET /api/analytics and GET /api/history APIs (with CSV export).
- Built the real-time Dashboard (totals, defect rate, latency, FPS, recent predictions, system status) auto-refreshing every 5 seconds.
- Implemented the System Monitor (CPU / memory / uptime) and a Human-in-the-Loop feedback endpoint for operator corrections.
- Generated the final evaluation: metrics, confusion matrix, ROC curve, and prediction-distribution charts.

6. Model Evaluation Results

Overall Performance (200-image validation set)

Metric	Value
Accuracy	67.0%
AUC (ROC)	0.7652
Validation Loss	0.5867
Avg. Inference Time	~280–620 ms (CPU)

Classification Report

Class	Precision	Recall	F1-Score	Support
Defect	0.77	0.48	0.59	100
Pass	0.62	0.86	0.72	100

Note: 67% accuracy is an expected baseline for highly randomized, procedurally-generated synthetic noise. The same architecture applied to a real-world dataset (e.g., the Kaggle PCB Defect Dataset) typically achieves 95%+ accuracy.

7. Conclusion

Over four weeks, the VisionSpec QC project delivered a complete, production-style AI pipeline for industrial visual inspection — from synthetic data generation and transfer-learning model training to Explainable AI, real-time APIs, and a full React dashboard. The architecture is modular, scalable, and ready to ingest real factory-floor data to reach state-of-the-art accuracy, complete with enterprise features such as Human-in-the-Loop logging and Grad-CAM diagnostics.

8. Future Enhancements

- Train on real-world PCB datasets to push accuracy beyond 95%.
- Multi-class defect categorization (short circuit, burn, missing solder, scratch).
- Dockerization and cloud deployment (Render backend + Vercel frontend).
- Edge-AI mode using TFLite / TensorRT for ultra-fast on-device inference.
- Auto-generated PDF QA reports and Slack / email defect-streak alerts.